The Relational Database Model

Overview

We now have a conceptual model for Monash Software, it is time to move to the second stage and map this to a logical model

For our unit this will involve mapping to the Relational Model in preparation for implementation in a RDBMS

- Relational Model
- Relational Algebra

The Relational Model

- Introduced by CODD in 1970 - the fundamental basis for the relational DBMS
- Basic structure is the mathematical concept of a RELATION mapped to the 'concept' of a table (tabular representation of relation)
    - Relation - abstract object
    - Table - pictorial representation
    - Storage structure - "real thing" - eg. isam file of 1's and 0's
- Relational Model Terminology
    - DOMAIN - set of atomic (indivisible) values
        - specify
            - name
            - data type
            - data format
- Examples:
    - customer_number domain - 5 character string of the form xxxdd
    - name domain - 20 character string
    - address domain - 30 character string containing street, town & postcode
    - credit_limit domain - money in the range $1,000 to $99,999

A Relation

- A relation consists of two parts
  - heading
  - body
- **Relation Heading:**
  - Also called Relational **Schema** consists of a fixed set of attributes
    - R (A1,A2,.......An)
      - R = relation name, Ai = attribute i
  - Each attribute corresponds to one underlying domain:
    - Customer relation heading:
      - CUSTOMER (custno, custname, custadd, credlimit)
        - » dom(custno) = customer_number
        - » dom(custname) = name
        - » dom(custadd) = address
        - » dom(credlimit) = credit_limit

custno	custname	custadd	credlimit

Relation Body

- **Relation Body**
  - Also called Relation Instance (state of the relation at any point in time)
    - r(R) = {t1, t2, t3, ..., tm}
    - consists of a time-varying set of n-tuples
      - Relation R consists of tuples t1, t2, t3 ., tm
      - m = number of tuples = **relation cardinality**
    - each n-tuple is an ordered list of n values
    - t = < v1, v2, ...vn>
      - n = number of values in tuple (no of attributes) = **relation degree**
  - In the tabular representation:
    - Relation heading ⇨ column headings
    - Relation body ⇨ set of data rows

custno	custname	custadd	credlimit
SMI13	SMITH	Wide Rd, Clayton, 3168	2000
JON44	JONES	Narrow St, Clayton, 3168	10000
BRO23	BROWN	Here Rd, Clayton, 3168	10000

Relation Properties

- No duplicate tuples
  - by definition sets do not contain duplicate elements
    - hence tuples must be unique
- Tuples are unordered within a relation
  - by definition sets are not ordered
    - hence tuples can only be accessed by content
- No ordering of attributes within a tuple
  - by definition sets are not ordered

Relation Properties cont'd

- Tuple values are atomic - cannot be divided

  - EMPLOYEE (eid, ename, departno, dependants)

    – not allowed: dependants (depname, depage) multivalued

  – hence no multivalued (repeating) attributes allowed, called the first normal form rule

- COMPARE with tabular representation

  - normally nothing to prevent duplicate rows

  - rows are ordered

  - columns are ordered

  – tables and relations are not the same 'thing'

Functional Dependency

- **Functional Dependency**:
  - A set of attributes A functionally determines an attribute B if, and only if, for each A value, there is exactly one value of B in the relation. It is denoted as A → B (A determines B, or B depends on A)
    - order_no → order_date
    - prod_no → prod_desc
    - order_no, prod_no → qty_ordered

ORDERNO	ORDERDATE
10	01/MAY/19
11	02/MAY/19
12	03/MAY/19
13	04/MAY/19
14	04/MAY/19
15	05/MAY/19
16	06/MAY/19

ORDERNO	PRODNO	QTYORDERED	LINEPRICE
10	101	1	11.98
11	101	1	11.98
11	103	2	123.58
12	104	10	479.8
13	105	2	140.36
14	106	1	31.99
15	107	3	116.73

PRODNO	PRODDESC	PRODUNITPRICE
101	Salmon – Smoked, Sliced	11.98
102	Broccolinni – Gaylan, Chinese	80.75
103	Pasta – Lasagne, Fresh	61.79
104	Melon – Cantaloupe	47.98
105	Wine – Peller Estates Late	70.18
106	Peas – Pigeon, Dry	31.99
107	Pumpkin – Seed	38.91

MONASH University

Relational Model Keys

- A **superkey** of a relation R is an attribute or set of attributes which exhibits only the uniqueness property

  – No two tuples of R have the same value for the superkey (Uniqueness property)

  – t1[superkey] ≠ t2[superkey]

- A **candidate key** CK of a relation R is an attribute or set of attributes which exhibits the following properties:

  – Uniqueness property (as above), *and*

  – No proper subset of CK has the uniqueness property (Minimality or Irreducibility property) ie. a minimal superkey

- One candidate key is chosen to be the **primary key** (PK) of a relation. Remaining candidate keys are termed alternate keys (AK).

Many possible superkeys

Potentially many possible candidate keys

Only ONE primary key (may be composed of many attributes - a composite primary key)

MONASH University

Selection of a Primary key

- **A primary key must be chosen considering the data that *may be added to the table in the future***
  - Names, dates of birth etc are rarely unique and as such are not a good option

  - PK should be free of 'extra' semantic meaning and security compliant, preferably a single attribute, preferably numeric (see Table 5.3 Coronel & Morris)

  - Natural vs Surrogate primary key

    - PATIENT_TREATMENT (patient_id, physician_id, treatment_code, pt_date, pt_time, pt_result)
      - Superkey
      - CK
      - PK
      - Issues with PK?

TABLE 5.3

DESIRABLE PRIMARY KEY CHARACTERISTICS

PK CHARACTERISTIC	RATIONALE
Unique values	The PK must uniquely identify each entity instance. A primary key must be able to guarantee unique values. It cannot contain nulls.
Nonintelligent	The PK should not have embedded semantic meaning other than to uniquely identify each entity instance. An attribute with embedded semantic meaning is probably better used as a descriptive characteristic of the entity than as an identifier. For example, a student ID of 65097 would be preferred over Smith, Martha T. as a primary key identifier.
No change over time	If an attribute has semantic meaning, it might be subject to updates, which is why names do not make good primary keys. If Vickie Smith is the primary key, what happens if she changes her name when she gets married? If a primary key is subject to change, the foreign key values must be updated, thus adding to the database work load. Furthermore, changing a primary key value means that you are basically changing the identity of an entity. In short, the PK should be permanent and unchangeable.
Preferably single-attribute	A primary key should have the minimum number of attributes possible (irreducible). Single-attribute primary keys are desirable but not required. Single-attribute primary keys simplify the implementation of foreign keys. Having multiple-attribute primary keys can cause primary keys of related entities to grow through the possible addition of many attributes, thus adding to the database workload and making (application) coding more cumbersome.
Preferably numeric	Unique values can be better managed when they are numeric, because the database can use internal routines to implement a counter-style attribute that automatically increments values with the addition of each new row. In fact, most database systems include the ability to use special constructs, such as Autonumber in Microsoft Access, sequence in Oracle, or uniqueidentifier in MS SQL Server to support self-incrementing primary key attributes.
Security-compliant	The selected primary key must not be composed of any attribute(s) that might be considered a security risk or violation. For example, using a Social Security number as a PK in an EMPLOYEE table is not a good idea.

Null in the Relational Model

- NULL is NOT a value - is a representation of the fact that there is *NO VALUE*
- Reasons for a NULL:
  - VALUE NOT APPLICABLE -
    - EMP relation - empno, deptno, salary, commission
      - commission only applies to staff in sales dept
  - VALUE UNKNOWN -
    - Joe's salary is NULL, Joe's salary is currently unknown
  - VALUE DOES NOT EXIST -
    - Tax File Number - is applicable to all employees but Joe may not have a number at this time
  - VALUE UNDEFINED -
    - Certain items explicitly undefined eg. divide by zero
      - Columns Number_of_payments, Total_payments
      - Column Average_payment_made
      - If Number_of_payments = 0 => Average undefined

MONASH University

Writing Relations

- Relations may be represented using the following notation:
  - RELATION_NAME (attribute1, attribute2,…)

- The primary key is underlined.

- Example:
  - STAFF (<u>staffid</u>, surname, initials, address, phone)

Relational Database

- A relational database is a collection of normalised relations.

- Normalisation is part of the design phase of the database and will be discussed in a later lecture.

Example relational database:

ORDER (<u>order_id</u>, orderdate,)

ORDER_LINE (<u>order_id</u>, <u>product_id</u>, quantity)

PRODUCT (<u>product_id</u>, description, unit_price)

Foreign Key (FK)

- FK: An attribute/s in a relation that exists in the same, or another relation as a Primary Key.

- Referential Integrity

  - **A Foreign key value must either *match* the *full primary key* in a relation or be NULL**.

- The pairing of PK and FK creates relationships (**logical** connections) between tables when implemented in a RDBMS. Hence the abstraction away from the underlying storage model.

Table Name: EMP					Table Name: DEPT		
EMPNO	ENAME	JOB	DEPTNO		DEPTNO	DNAME	LOC
7839	KING	PRESIDENT	10		10	ACCOUNTING	NEW YORK
7698	BLAKE	MANAGER	30		20	RESEARCH	DALLAS
7782	CLARK	MANAGER	10		30	SALES	CHICAGO
7566	JONES	MANAGER	20		40	OPERATIONS	BOSTON

Primary key   Foreign key   Primary key

Data Integrity

- Entity integrity
  - Primary key value must not be NULL.
    - No duplicate tuple property then ensures that each primary key must be unique
- Referential integrity
  - The values of FK must either match a value of a full PK in the related relation or be NULL.
- Column/Domain integrity
  - All values in a given column must come from the same domain (the same data type and range).

Relational DMLs

- Relational Calculus
- Relational Algebra
- Transform Oriented Languages (e.g. SQL)
- Graphical Languages
- Exhibit the "closure" property - queries on relations produce relations

Relational Calculus

- Based on mathematical logic.

- Non-procedural.

- Primarily of theoretical importance.

- May be used as a yardstick for measuring the power of other relational languages ("relational completeness").

- Operators may be applied to any number of relations.

MONASH
University

RELATIONAL ALGEBRA

Manipulation of relational data

MONASH
University

Relational Algebra

- Relationally complete.
- Procedural.
- Operators only apply to at most two relations at a time.
- 8 basic operations:
  - single relation: selection, projection
  - cartesian product, join
  - union
  - intersection
  - difference
  - division

Relational Operation PROJECT

$$\pi$$

PROJ_CODE	PROJ_MANAGER	PROJ_BID_PRICE
21-5Z	Holly B. Parker	$16,833,460.00
25-2D	...	$13,500,000.00
25-5A	George F. Dorts	$32,512,420.00
25-9T	Holly B. Parker	$21,563,234.00
27-4Q	George F. Dorts	$10,314,545.00
29-2D	Holly B. Parker	$25,559,999.00
31-7P	William K. Moor	$56,850,000.00

Relational Operation SELECT

$$\sigma$$

PROJECT_CODE	PROJECT_MANAGER	PROJECT_BID_PRICE
21-5Z	Holly B. Parker	$16,833,460.00
25-2D	Jane D. Grant	$12,500,000.00
25-5A	George F. Dorts	$32,512,420.00
25-9T	Holly B. Parker	$21,563,234.00
27-4Q	George F. Dorts	$10,314,545.00
29-2D	Holly B. Parker	$25,559,999.00
31-7P	William K. Moor	$56,850,000.00

Relational Operation Multiple Actions

2

PROJECT_CODE	PROJECT_MANAGER	PROJECT_BID_PRICE
21-5T	Holly B. Parker	$15,254,000.00
25-2D	Jane D. Grant	$12,500,000.00
1 25-5A	George F. Dorts	$32,512,420.00
25-9T	Holly B. Parker	$21,563,234.00
27-4Q	George F. Dorts	$10,314,545.00
29-2D	Holly B. Parker	$25,559,999.00
31-7P	William K. Moor	$56,850,000.00

$$\text{Result} = \pi_{\text{project\_manager}}(\sigma_{\text{project\_code}=25\text{-}5A}\text{PROJECT})$$

SQL vs Relational Algebra in the Database

Assignment Project Exam Help

SQL Queries

https://powcoder.com

```
SELECT project_manager
FROM project
WHERE project_code='25-5A';
```

Query Processor

Add WeChat powcoder

Relational Algebra Operations

$$\pi_{\text{project\_manager}}(\sigma_{\text{project\_code}=25\text{-}5A}\text{PROJECT})$$

DBMS
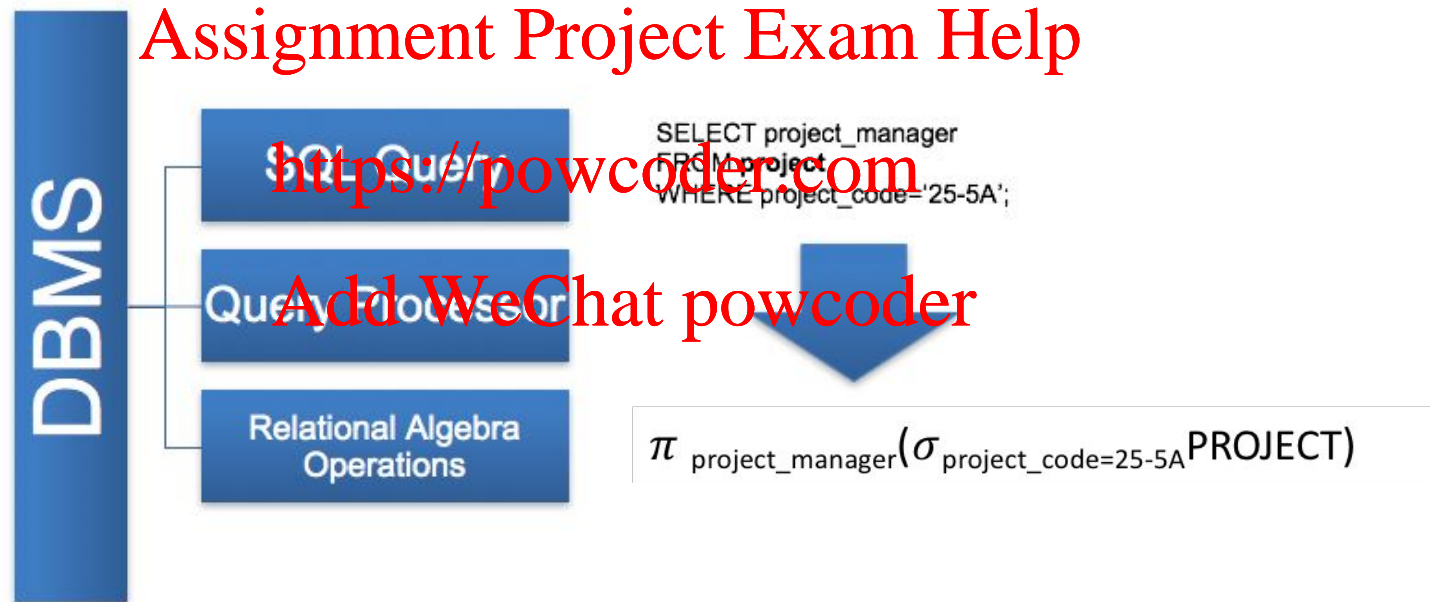
JOIN

- Join operator used to combine data from two or more relations, based on a common attribute or attributes.
- Different types:
  - theta-join
  - equi-join
  - natural join
  - outer join

Assignment Project Exam Help

https://powcoder.com

Add WeChat powcoder

THETA JOIN (Generalised join)

$$(\text{Relation\_1}) \bowtie_F (\text{Relation\_2})$$

- *F* is a predicate (i.e. truth-valued function) which is of the form $\text{Relation\_1}.a_i$ θ $\text{Relation2}.b_i$
  - CUSTOMER.cust_no θ ORDER.cust_no
- θ is one of the standard arithmetic comparison operators, i.e. <, ≤, =, ≥, >
- Most commonly, θ is equals (=), but can be *any* of the operators
  - EMPLOYEE.emp_sal > SALARYSCALE.step_5

NATURAL JOIN

STUDENT			MARK		
ID	Name		ID	Subj	Marks
1	Alice		1	1004	95
2	Bob		2	1045	55
			1	1045	90

Step 1: STUDENT X MARK
Step 2: delete rows where IDs do not match (select =)

STUDENT.ID	Name	MARK.ID	Subj	Marks
1	Alice	1	1004	95
1	Alice	2	1045	55
1	Alice	1	1045	90
2	Bob	1	1004	95
2	Bob	2	1045	55
2	Bob	1	1045	90

MONASH University

37

NATURAL JOIN

STUDENT			MARK		
ID	Name		ID	Subj	Marks
1	Alice		1	1004	95
2	Bob		2	1045	55
				1044	90

Step 1: STUDENT X MARK
Step 2: delete rows where IDs do not match (select =)
Step 3: delete duplicate columns (project away)

STUDENT.ID	Name	MARK.ID	Subj	Marks
1	Alice	1	1004	95
1	Alice	1	1045	90
2	Bob	2	1045	55

NATURAL JOIN

STUDENT

ID	Name
1	Alice
2	Bob

MARK

ID	Subj	Marks
1	1004	95
2	1045	55
1	1045	90

Step 1: STUDENT X MARK
Step 2: delete rows where IDs do not match (select =)
Step 3: delete duplicate columns (project away)

ID	Name	Subj	Marks
1	Alice	1004	95
1	Alice	1045	90
2	Bob	1045	55

A natural join of STUDENT and MARK

MONASH University

39

OUTER JOIN

STUDENT

ID	Name
1	Alice
2	Bob
3	Chris

MARK

ID	Subj	Marks
1	1004	95
2	1045	55
1	1045	90
4	1004	100

⋈

No information for Chris (no mark, e.g. just enrolled), and the student with ID 4 (no student, e.g. quit uni)

ID	Name	Subj	Marks
1	Alice	1004	95
1	Alice	1045	90
2	Bob	1045	55

A natural join of STUDENT and MARK

FULL OUTER JOIN

STUDENT

ID	Name
1	Alice
2	Bob
3	Chris

MARK

ID	Subj	Marks
1	1004	95
2	1045	55
1	1045	90
4	1004	100

⋈

Get (incomplete) information of both Chris and student with ID 4

ID	Name	Subj	Marks
1	Alice	1004	95
1	Alice	1045	90
2	Bob	1045	55
3	Chris	Null	Null
4	Null	1004	100

A full outer join of STUDENT and MARK

42

LEFT OUTER JOIN

STUDENT			MARK		
ID	Name		ID	Subj	Marks
1	Alice		1	1004	95
2	Bob		2	1045	55
3	Chris		1	1045	90
			4	1004	100

← Get (incomplete) information of only Chris

ID	Name	Subj	Marks
1	Alice	1004	95
1	Alice	1045	90
2	Bob	1045	55
3	Chris	Null	Null

A left outer join of STUDENT and MARK
Memory aid: Chris is on the ← LEFT of the nulls.

MONASH University

RIGHT OUTER JOIN

STUDENT

ID	Name
1	Alice
2	Bob
3	Chris

MARK

ID	Subj	Marks
1	1004	95
2	1045	55
1	1045	90
4	1004	100

⋈

Get (incomplete) information of the student with ID 4 →

ID	Name	Subj	Marks
1	Alice	1004	95
1	Alice	1045	90
2	Bob	1045	55
4	Null	1004	100

A right outer join of STUDENT and MARK.
Memory aid: the marks data is on the RIGHT → of the null.